The background features a light blue sky with a large yellow sun in the top right corner and several white, fluffy clouds. The bottom of the slide shows a green landscape with rolling hills, stylized green trees with brown trunks, and small pink flowers.

Using segmentation to build a capability-based single address space operating system

Wuyang Chung
wy-chung@outlook.com

You can get the slides from <https://github.com/wy-chung/segmentation>

Outline

- Single Address Space Operating System
- Capability-based Addressing
- Segmentation
- Benefits of Segmentation
- Compatibility with Old Programs
- Conclusion



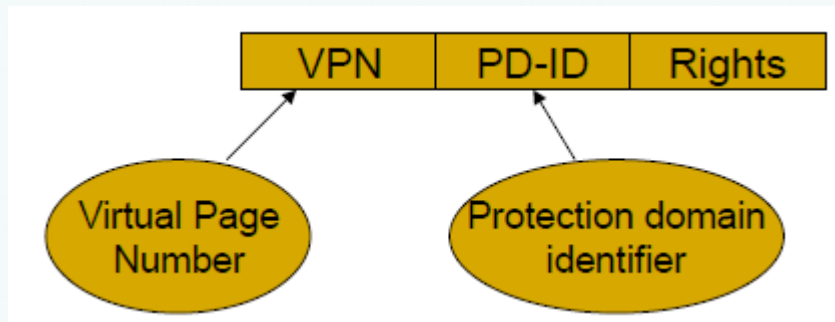
Single Address Space OS

- Problems with multiple address space OS
 - Pointers cannot be used in shared buffer
 - On CPUs without ASID
 - Need to flush TLB when doing address space switch
 - On CPUs with ASID
 - Duplications of translation information in TLB
- Advantages of single address space OS
 - Pointers can be used in shared buffer
 - Lower context switch overhead
- Challenge of single address space OS
 - Protection



Multiple Protection Domains on a Single Address Space

- Domain page model
 - PLB (Protection Lookaside Buffer)



Architecture support for single address space operating systems, ACM SIGPLAN, 1992.

https://cseweb.ucsd.edu/classes/fa11/cse240A-a/Slides1/06_SAOS.pdf

- Page group model
 - There is an AID (access identifier) field in page table entry
 - A process can access a page if any of the PIDs is equal to the page AID
 - The CPU has several PID (protection identifier) registers
- Capability-based addressing



Capability-based Addressing


capability permissions unique object ID



capability-based
addressing permissions unique segment ID

- The most important thing in a capability system is to prevent capability from being forged
- Two approaches
 - Tagged approach
 - e.g. CHERI
 - Partitioned approach
 - e.g. Plessey 250

Partitioned approach




Capability-based Addressing

capability permissions unique object ID



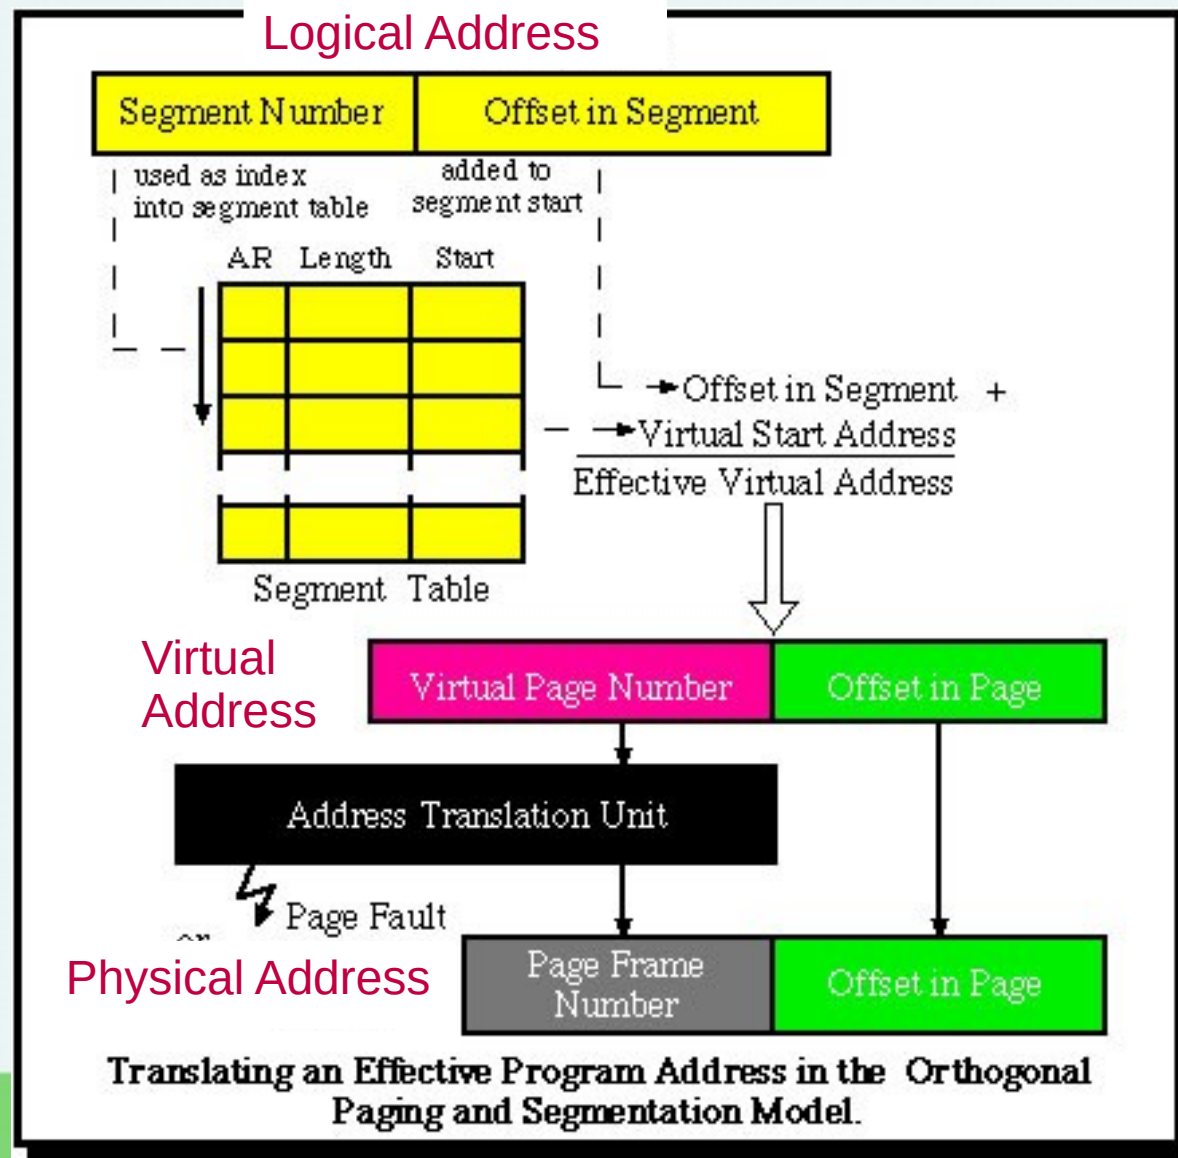
capability-based
addressing permissions unique segment ID

- The most important thing in a capability system is to prevent capability from being forged
- Two approaches
 - Tagged approach
 - e.g. CHERI
 - Partitioned approach
 - e.g. Plessey 250



Orthogonal Segmentation and Paging

<https://www.monads-security.org/orthogonal-segmentation-and-paging.html>



Segmentation

capability descriptor

tag index_c

code
data
stack

p0 GPR

p1 GPR

code
data
stack

capability table
(C-list)

tag	perms	index _s

ctr base length

tag	perms	index _s

kernel

segment metadata table

tag	base	size	attrs

str base length

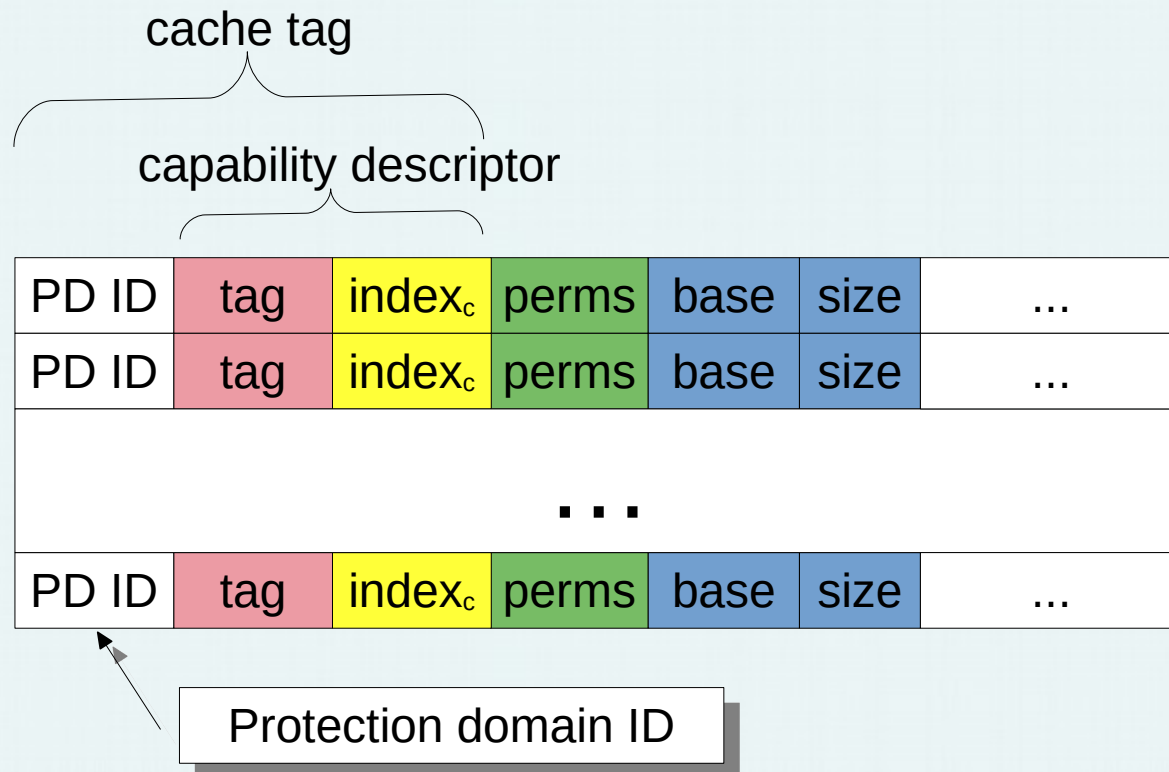
capability

permissions unique segment ID

Tag must have at least 4 bits

Segmentation Hardware

- SLB (Segment Lookaside Buffer)
 - Cache permissions and segment descriptors



Segment Types

- Segmentation translates logical address to linear address
 - There are 3 types of linear address
 - Virtual, physical and I/O address
- Segment types
 - Virtual segment
 - Linear address is virtual address
 - Physical segment
 - Linear address is physical address
 - I/O segment
 - Linear address is I/O address

logical address is
{capability descriptor, offset within the segment}



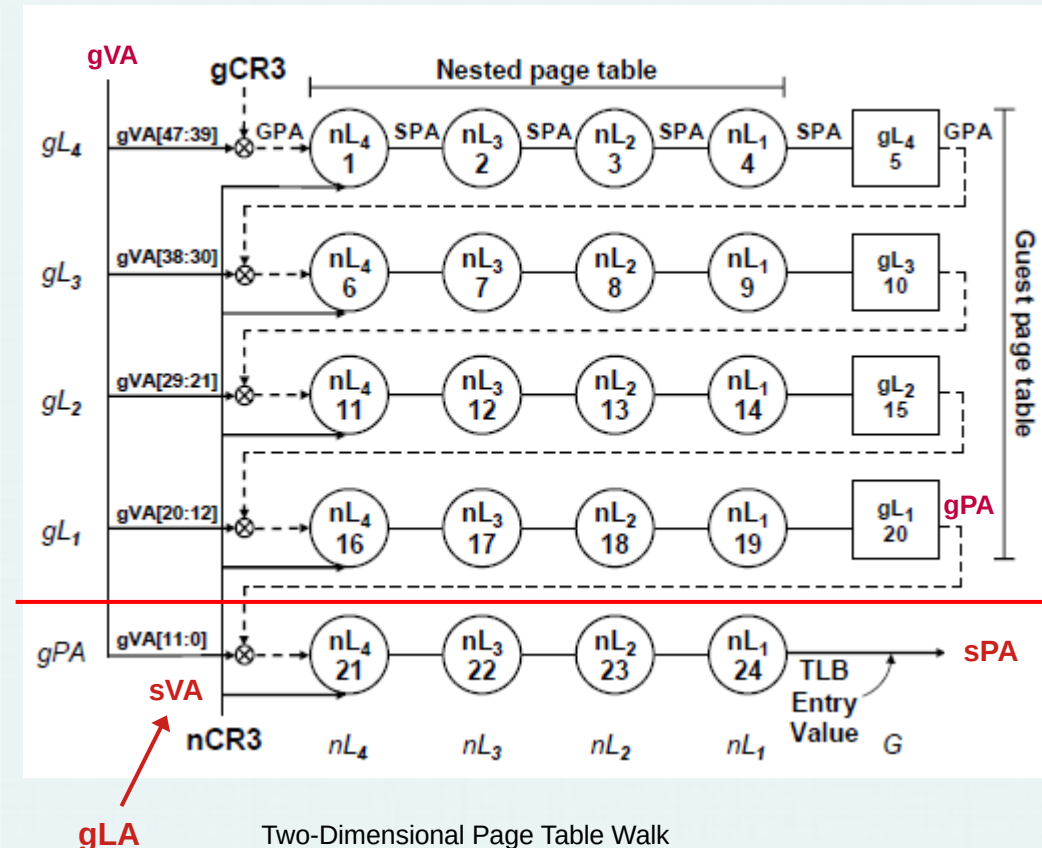
Benefits of Virtual Segment

- Fast Segment move
 - Don't need to copy the segment data, only need to copy the virtual to physical mapping
 - Don't need to fix any pointers in the segment after the segment is moved



Simplification of page table walk in Guest Virtual Machine

- On a guest machine, hardware needs to do two virtual address translations for TLB miss
 - Guest virtual address to guest physical address
 - System virtual address to system physical address
- With virtual segment only one virtual address translation is needed
 - Guest logical address can be translated directly to system virtual address



Two-Dimensional Page Table Walk

Accelerating two-dimensional page walks for virtualized systems, acm ASPLOS, 2008

logical address is
{capability descriptor, offset within the segment}




Benefits of Physical Segment

- Software-managed TLB
 - Must guarantee no TLB miss exceptions when the kernel is running the TLB miss handling routine
- Direct segment
 - Arkaprava Basu, Jayneel Gandhi, Jichuan Chang, Mark D. Hill, Michael M. Swift. Efficient virtual memory for big memory servers. In Proc. ISCA, 2013.
 - TLB misses consume up to 51% of execution cycles for big memory servers
- IOMMU simplification
 - Simple authorization
 - No page translation is needed
 - IOMMU is simply a cache for physical segment metadata



Simplification of IOMMU

- How the IOMMU protects main memory from hardware device?
 1. Device driver allocates a physical segment and sends the capability of that segment to hardware device
 2. Device sends the **capability**, offset and data length to IOMMU
 3. IOMMU verifies that the capability is valid and it points to a **physical segment**
 4. IOMMU then checks that the offset and length is within segment size
 5. IOMMU calculates the physical address by adding segment base address and the offset
 6. The calculated physical address is used to access system main memory



From Wikipedia



I/O Segment

- Purpose
 - Leave memory address space to memory only and move all the other stuff to I/O address space, such as CPU's CSRs, PCI configuration space, video frame buffer, boot ROM, etc.
- Benefits
 - No memory holes in memory address space
 - Segment is more fine-grained than page
 - For a large I/O buffer, only one SLB entry is needed
 - The same load instruction can be used to load from either I/O or memory address space
 - Even when a device driver runs at the user level, it can access its hardware device without kernel intervention



Simplification of Shared Library

- No need for PIC (Position Independent Code), GOT (Global Offset Table) and PLT (Procedure Linkage Table)
 - Shared library has its own code and data segments and segment always starts with offset 0
- Shared libraries can be partially linked
 - The offset part of the library's global variables and functions can be statically linked at compile time
 - The capability descriptor of the library's global variables and functions are linked at program load time



Cross-process Call

- Cross-process call has many names: cross-domain call, migrating thread model, protected control transfer, passive object model
- Traditionally thread is confined in two protection domains
 - When a thread is in the kernel, it can only return back to its creator process
- Cross-process Call
 - When a thread is in the kernel, it is allowed to upcall into other process
- Benefits
 - Encourage modularity and improve security
 - Resource accounting can be made more accurate
 - Preferable for real-time systems

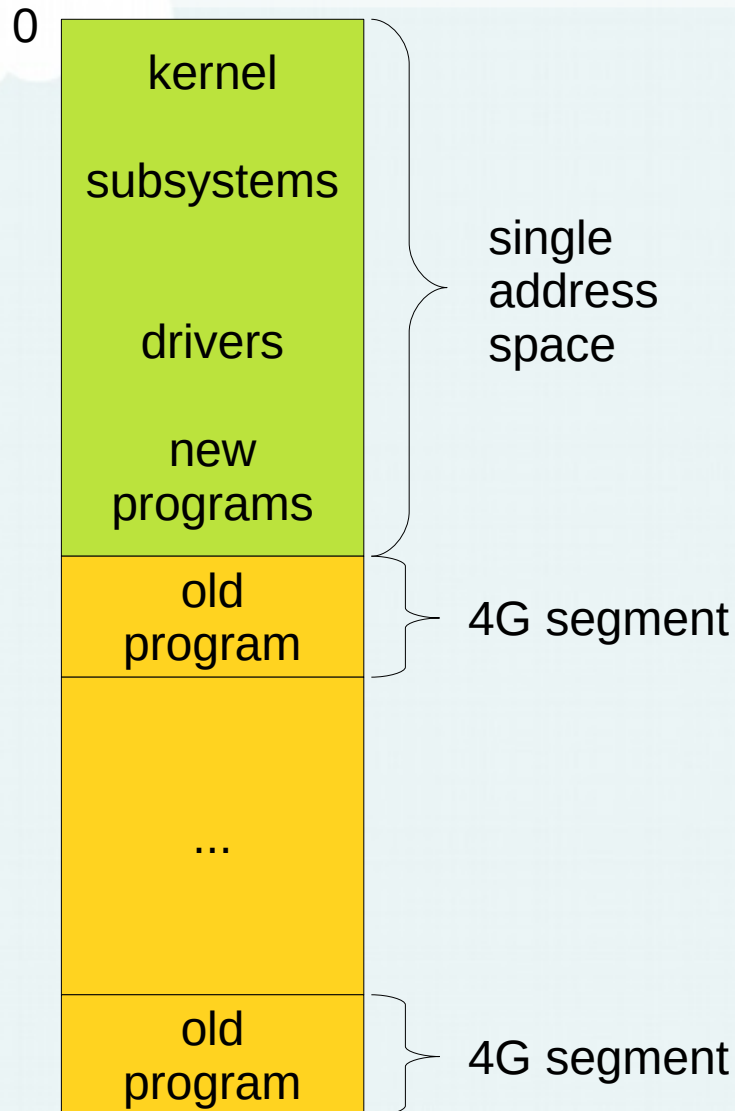


Simplification of Cross-process Call

- The same stack segment can be used in both client and server process if we can
 - Protect the server part of the stack from client
 - Move the stack capability from client process to server process
 - Protect the client part of the stack from server
 - Stack boundary pointer (sbp) register
 - A privileged register
 - CPU restricts that a thread can only access server part of the stack



Compatibility with Old Programs



- Ideal platform
 - 32-bit programs on 64-bit virtual address space
- The first half of the virtual address space is for kernel, subsystems, device drivers and single-address-space programs
- The second half of the virtual address space is reserved for old programs



Conclusion

- Lots of benefits we can get from single address space, capability and segmentation
- Changes
 - OS
 - Segment, Capability and single address space
 - Compiler
 - Far pointer, far function call
 - New shared library implementation
 - Hardware
 - Capability and segmentation architecture






The End

Thank you




Stack Growth Direction

- Stack should grow towards ∞ instead of 0
 - Can expand the stack segment when it reaches its maximum size



towards 0



towards ∞

