Using segmentation to build a capability-based single address space operating system

Wuyang Chung

wy-chung@outlook.com

You can see the video on YouTube.

Outline

- Single Address Space Operating System
- Capability-based Addressing
- Segmentation
- Benefits of Segmentation
- Compatibility with Old Programs
- Conclusion

Single Address Space OS

- Problems with multiple address space OS
  - Pointers cannot be used in shared buffer
  - On CPUs without ASID
    - Need to flush TLB when doing address space switch
  - On CPUs with ASID
    - Duplications of translation information in TLB
- Advantages of single address space OS
  - Pointers can be used in shared buffer
  - Lower context switch overhead
- Challenge of single address space OS
  - Protection

Multiple Protection Domains
on a Single Address Space

- Domain page model

  - PLB (Protection Lookaside Buffer)
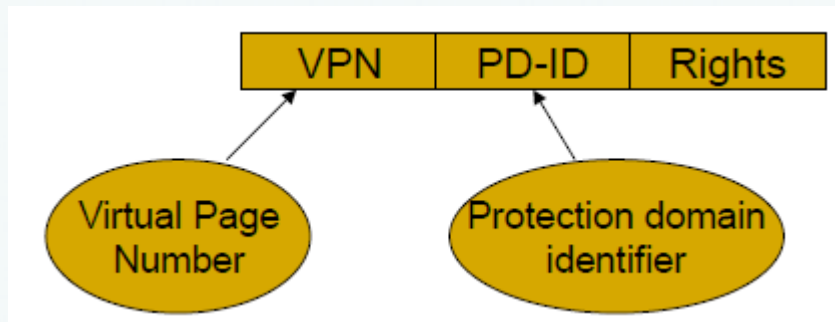
Architecture support for single address space operating systems, ACM SIGPLAN, 1992.

https://cseweb.ucsd.edu/classes/fa11/cse240A-a/Slides1/06_SAOS.pdf

- Page group model

  - There is an AID (access identifier) field in page table entry

  - A process can access a page if any of the PIDs is equal to the page AID

  - The CPU has several PID (protection identifier) registers
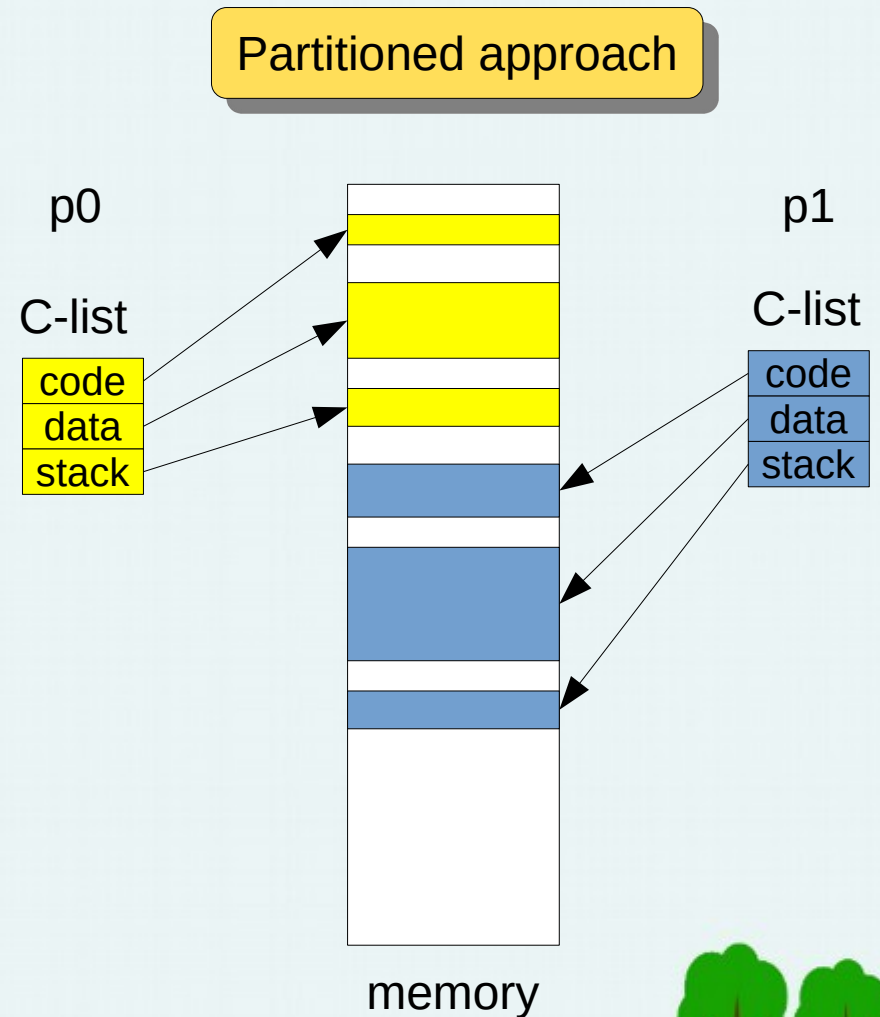
- Capability-based addressing

Capability-based Addressing

capability | permissions | unique object ID

↓

capability-based addressing | permissions | unique segment ID

Partitioned approach

- The most important thing in a capability system is to prevent capability from being forged

- Two approaches
  - Tagged approach
    - e.g. CHERI
  - Partitioned approach
    - e.g. Plessey 250

p0

C-list
| code |
| data |
| stack |

p1

C-list
| code |
| data |
| stack |

memory

Capability-based Addressing

capability | permissions | unique object ID |

↓

capability-based addressing | permissions | unique segment ID |

- The most important thing in a capability system is to prevent capability from being forged

- Two approaches
  - Tagged approach
    - e.g. CHERI
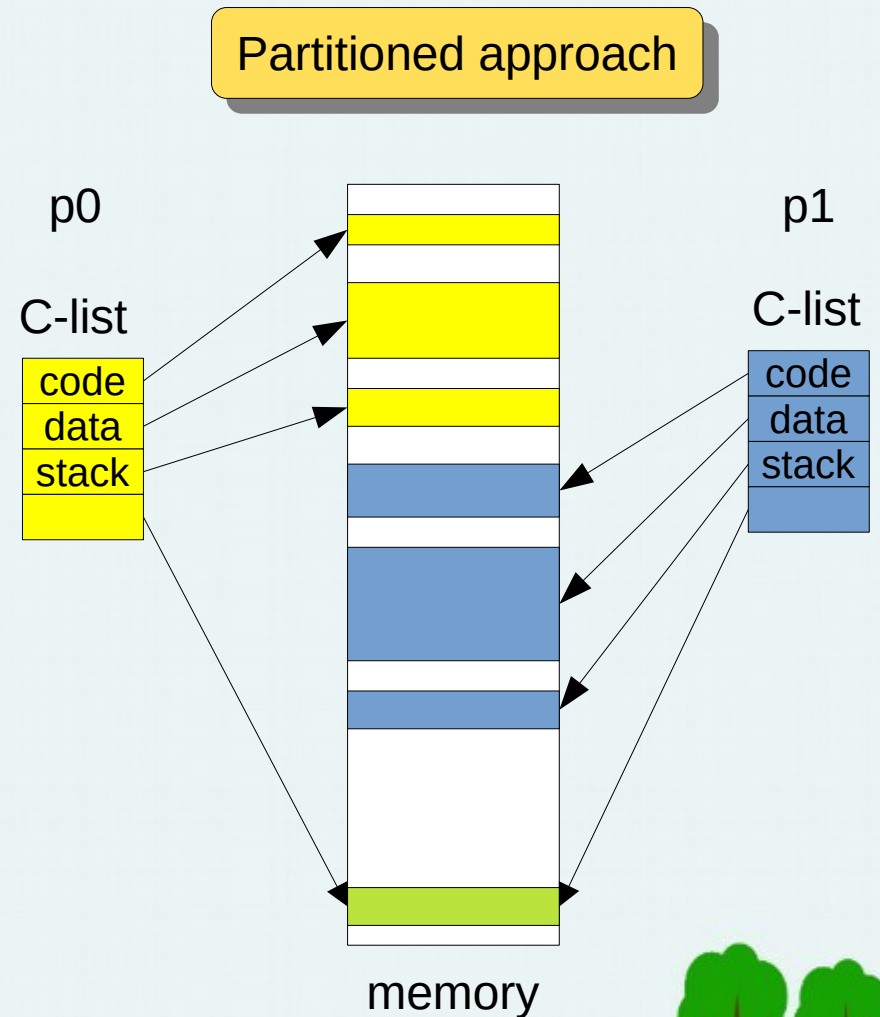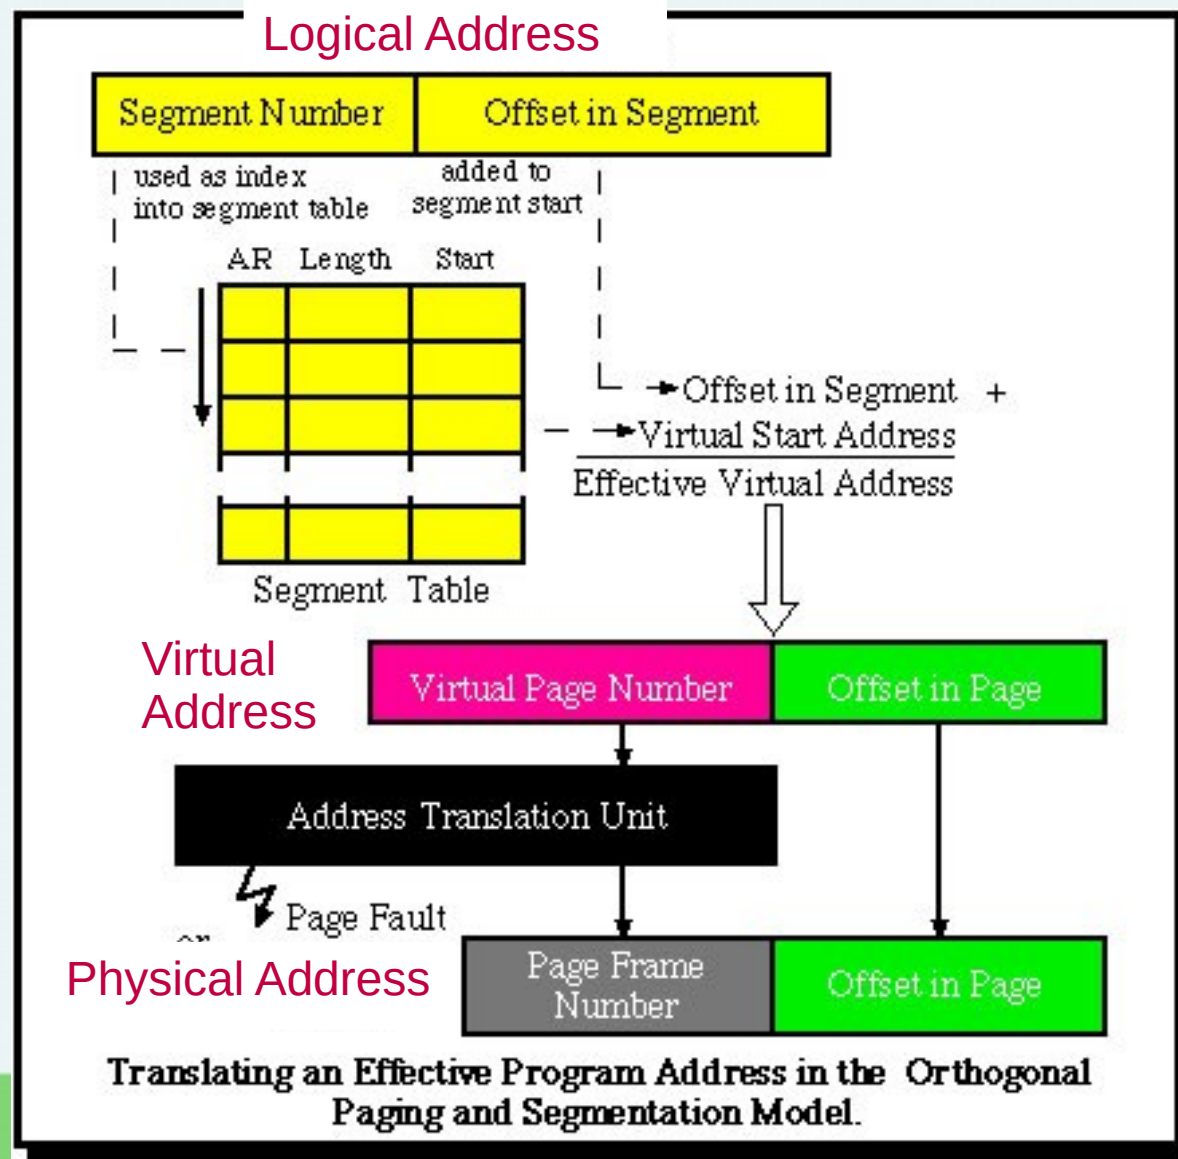 - Partitioned approach
    - e.g. Plessey 250

Partitioned approach

p0

C-list

| code |
| data |
| stack |

p1

C-list

| code |
| data |
| stack |

memory

Orthogonal Segmentation and Paging

https://www.monads-security.org/orthogonal-segmentation-and-paging.html

Logical Address

Virtual Address

Physical Address

Translating an Effective Program Address in the Orthogonal Paging and Segmentation Model.

7

Segmentation

capability descriptor

capability table (C-list)

segment metadata table

tag	index$_c$

	code	
	data	
	stack	

p0   GPR

p1   GPR

	code	
	data	
	stack	

tag	perms	index$_s$

ctr | base | length |

tag	perms	index$_s$

kernel

tag	base	size	attrs

str | base | length |

capability

permissions	unique segment ID

Tag must have at least 4 bits

8

Segmentation Hardware

- SLB (Segment Lookaside Buffer)
    - Cache permissions and segment descriptors

cache tag

capability descriptor

PD ID	tag	index$_c$	perms	base	size	...
PD ID	tag	index$_c$	perms	base	size	...
...						
PD ID	tag	index$_c$	perms	base	size	...

Protection domain ID
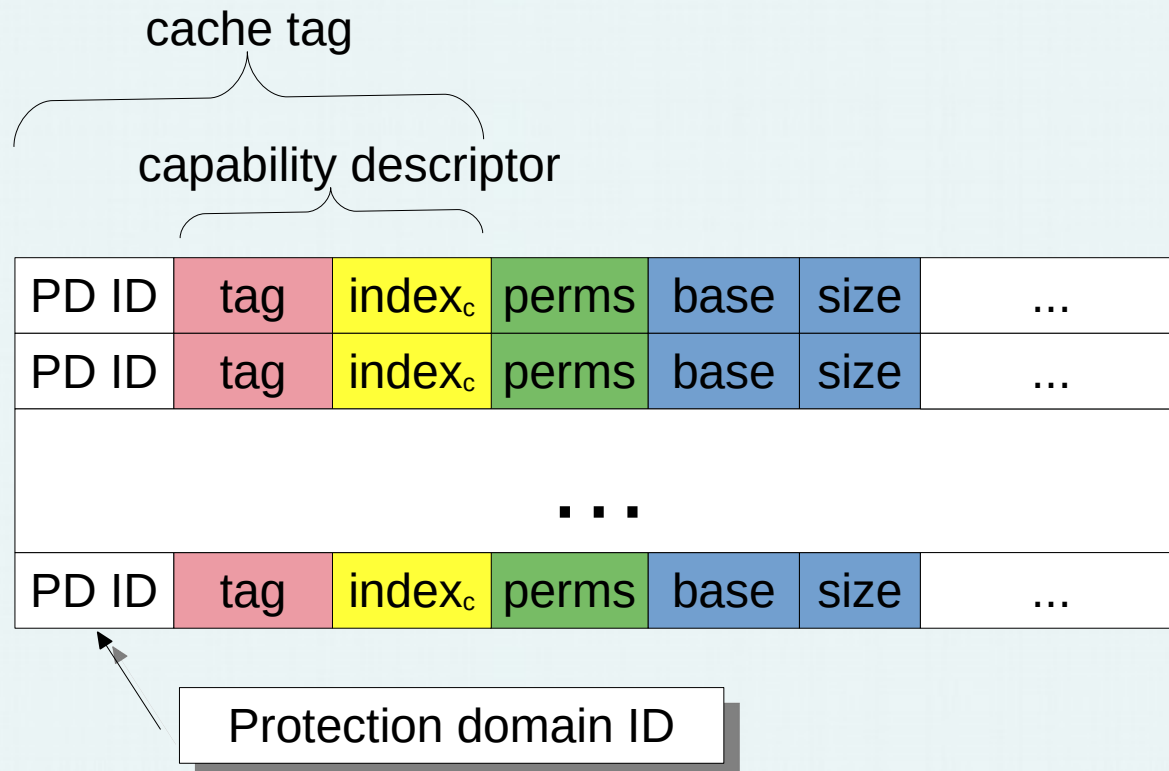
Segment Types

- Segmentation translates logical address to linear address
  - There are 3 types of linear address
    - Virtual, physical and I/O address
- Segment types

  > logical address is
  > {capability descriptor, offset within the segment}

  - Virtual segment
    - Linear address is virtual address
  - Physical segment
    - Linear address is physical address
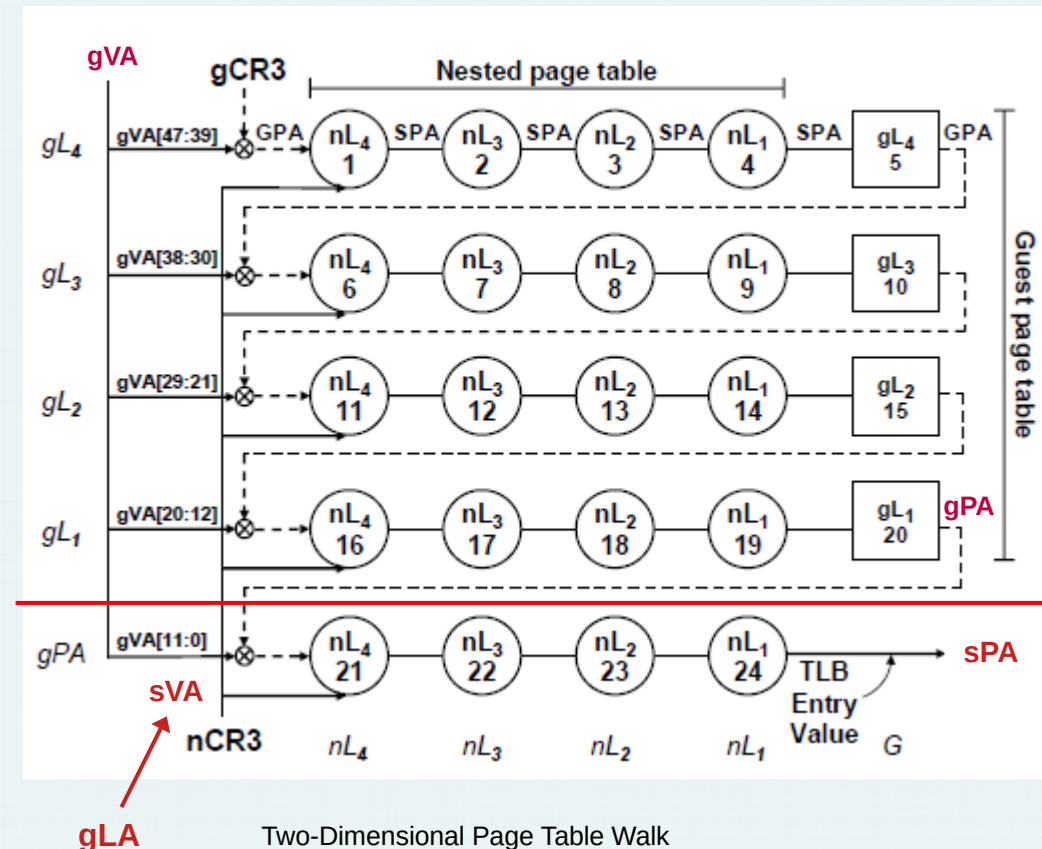 - I/O segment
    - Linear address is I/O address

Benefits of Virtual Segment

- Fast Segment move

    – Don't need to copy the segment data, only need to copy the virtual to physical mapping

    – Don't need to fix any pointers in the segment after the segment is moved

Simplification of page table walk in Guest Virtual Machine

- On a guest machine, hardware needs to do two virtual address translations for TLB miss

  – Guest virtual address to guest physical address

  – System virtual address to system physical address

- With virtual segment only one virtual address translation is needed

  – Guest logical address can be translated directly to system virtual address

Two-Dimensional Page Table Walk

Accelerating two-dimensional page walks for virtualized systems, acm ASPLOS, 2008

logical address is
{capability descriptor, offset within the segment}

Benefits of Physical Segment

- Software-manged TLB
  - Must guarantee no TLB miss exceptions when the kernel is running the TLB miss handling routine

- Direct segment
  - Arkaprava Basu, Jayneel Gandhi, Jichuan Chang, Mark D. Hill, Michael M. Swift. Efficient virtual memory for big memory servers. In Proc. ISCA, 2013.
    - TLB misses consume up to 51% of execution cycles for big memory servers

- IOMMU simplification
  - Simple authorization
  - No page translation is needed
  - IOMMU is simply a cache for physical segment metadata

Simplification of IOMMU

- How the IOMMU protects main memory from hardware device?

  1. Device driver allocates a physical segment and sends the capability of that segment to hardware device

  2. Device sends the **capability**, offset and data length to IOMMU

  3. IOMMU verifies that the capability is valid and it points to a **physical segment**

  4. IOMMU then checks that the offset and length is within segment size

  5. IOMMU calculates the physical address by adding segment base address and the offset

  6. The calculated physical address is used to access system main memory

Comparison of the I/O memory management unit (IOMMU) to the memory management unit (MMU).
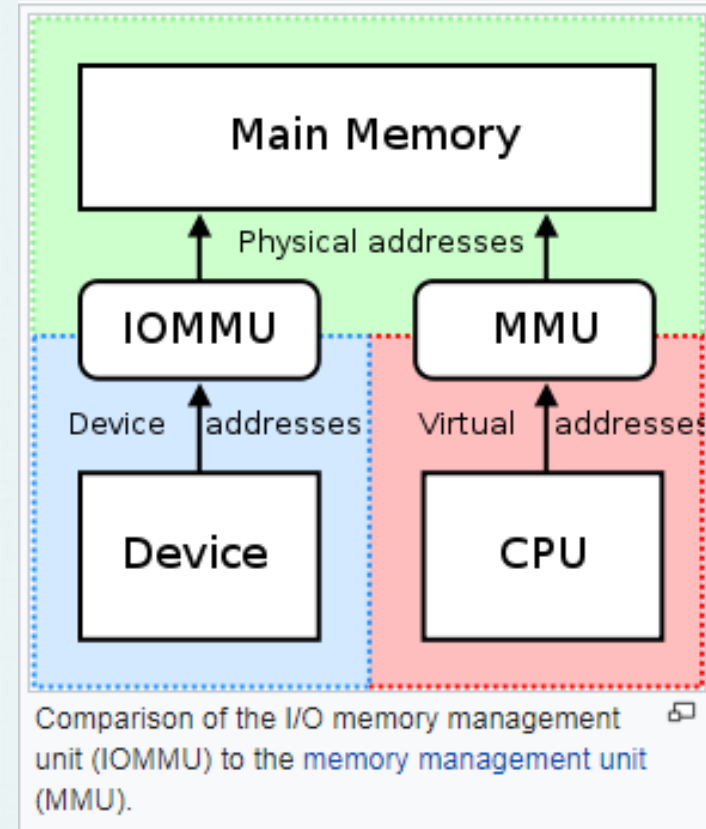
From Wikipedia

I/O Segment

- Purpose
  - Leave memory address space to memory only and move all the other stuff to I/O address space, such as CPU's CSRs, PCI configuration space, video frame buffer, boot ROM, etc.

- Benefits
  - No memory holes in memory address space
  - Segment is more fine-grained then page
  - For a large I/O buffer, only one SLB entry is needed
  - The same load instruction can be used to load from either I/O or memory address space
  - Even when a device driver runs at the user level, it can access its hardware device without kernel intervention

Simplification of Shared Library

- No need for PIC (Position Independent Code), GOT (Global Offset Table) and PLT (Procedure Linkage Table)

  - Shared library has its own code and data segments and segment always starts with offset 0

- Shared libraries can be partially linked

  - The offset part of the library's global variables and functions can be statically linked at compile time

  - The capability descriptor of the library's global variables and functions are linked at program load time

Cross-process Call

- Cross-process call has many names: cross-domain call, migrating thread model, protected control transfer, passive object model

- Traditionally thread is confined in two protection domains
  - When a thread is in the kernel, it can only return back to its creator process

- Cross-process Call
  - When a thread is in the kernel, it is allowed to upcall into other process

- Benefits
  - Encourage modularity and improve security
  - Resource accounting can be made more accurate
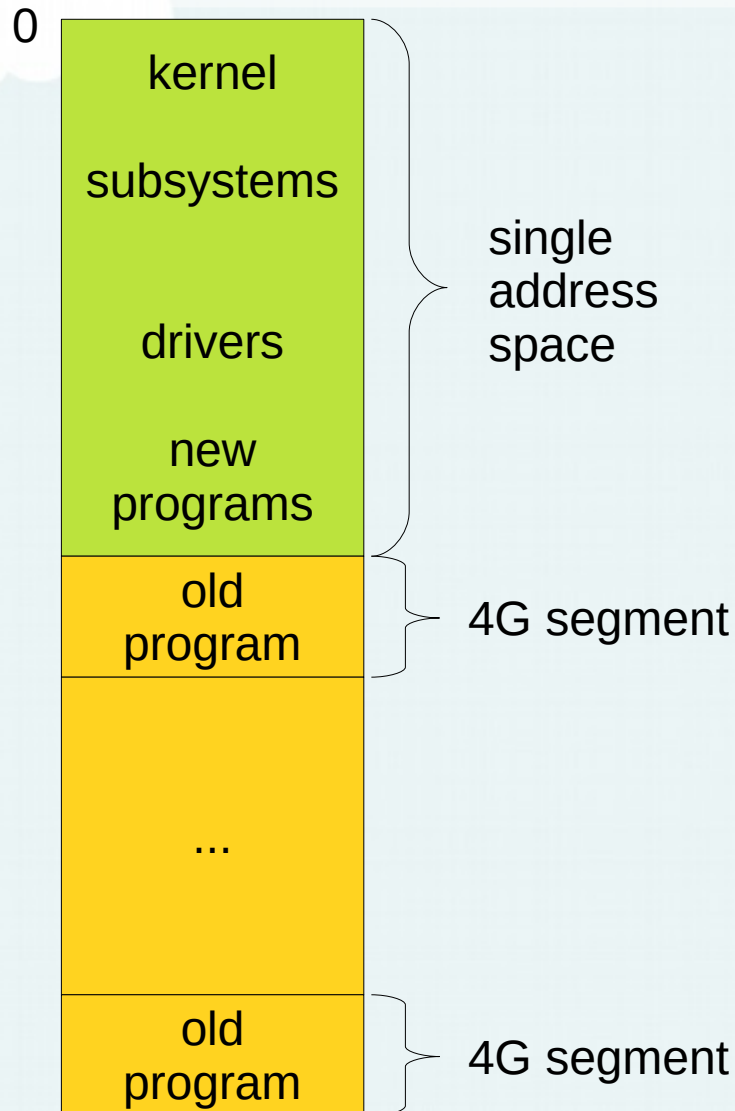 - Preferable for real-time systems

Simplification of Cross-process Call

- The same stack segment can be used in both client and server process if we can

  - Protect the server part of the stack from client

    - Move the stack capability from client process to server process

  - Protect the client part of the stack from server

    - Stack boundary pointer (sbp) register

      - A privileged register
      - CPU restricts that a thread can only access server part of the stack

Compatibility with Old Programs

```
0
┌──────────────┐
│   kernel     │
│              │
│ subsystems   │  ┐
│              │  ├ single
│              │  │ address
│  drivers     │  │ space
│              │  ┘
│    new       │
│ programs     │
├──────────────┤
│   old        │  ┐ 4G segment
│ program      │  ┘
│              │
│              │
│     ...      │
│              │
├──────────────┤
│   old        │  ┐ 4G segment
│ program      │  ┘
└──────────────┘
```

- Ideal platform
  - 32-bit programs on 64-bit virtual address space
- The first half of the virtual address space is for kernel, subsystems, device drivers and single-address-space programs
- The second half of the virtual address space is reserved for old programs

Conclusion

- Lots of benefits we can get from single address space, capability and segmentation

- Changes
  - OS
    - Segment, Capability and single address space
  - Compiler
    - Far pointer, far function call
    - New shared library implementation
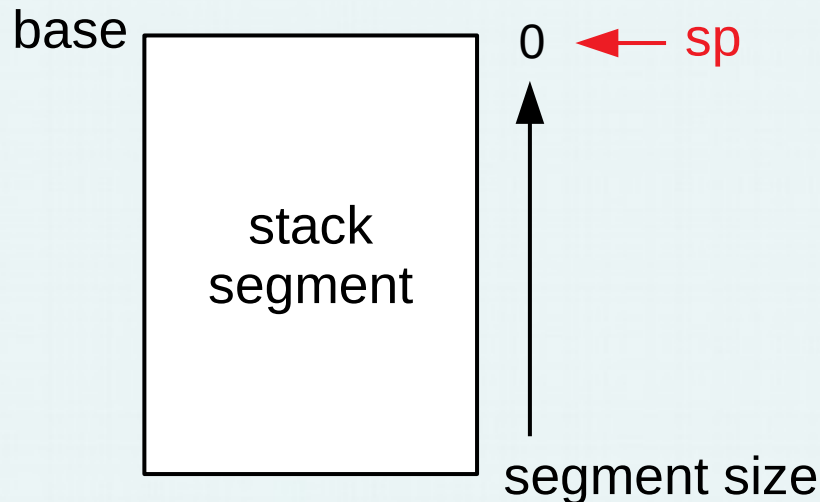 - Hardware
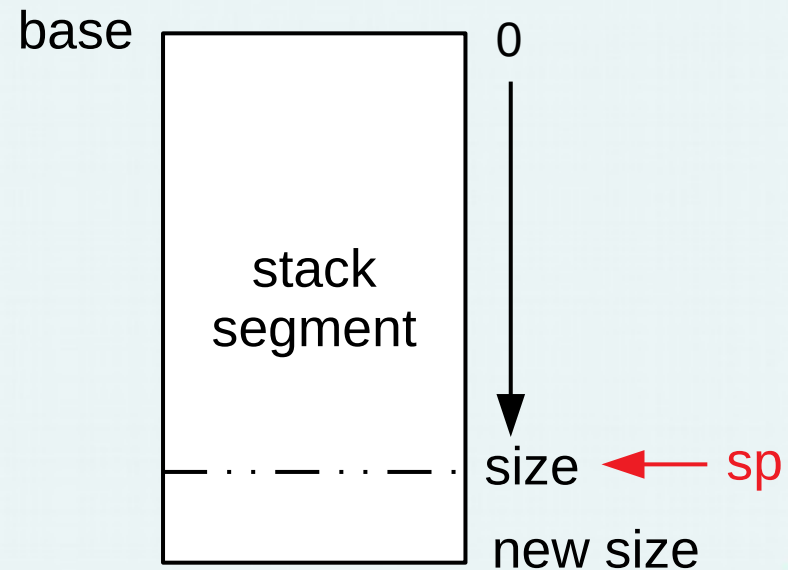   - Capability and segmentation architecture

The End

Thank you

Stack Growth Direction

- Stack should grow towards ∞ instead of 0
  - Can expand the stack segment when it reaches its maximum size

base   stack segment   0 ← sp

segment size

towards 0

base   stack segment   0

size ← sp

new size

towards ∞